

Introduction to CUDA

1 Our first GPU Program

- running Newton's method in complex arithmetic
- examining the CUDA Compute Capability

2 CUDA Program Structure

- steps to write code for the GPU
- code to compute complex roots
- the kernel function and main program
- a scalable programming model

MCS 572 Lecture 30
Introduction to Supercomputing
Jan Verschelde, 31 October 2016

Introduction to CUDA

1 Our first GPU Program

- running Newton's method in complex arithmetic
- examining the CUDA Compute Capability

2 CUDA Program Structure

- steps to write code for the GPU
- code to compute complex roots
- the kernel function and main program
- a scalable programming model

computing complex square roots

To compute \sqrt{c} for $c \in \mathbb{C}$,
we apply Newton's method on $x^2 - c = 0$:

$$x_0 := c, \quad x_{k+1} := x_k - \frac{x_k^2 - c}{2x_k}, \quad k = 0, 1, \dots$$

Five iterations suffice to obtain an accurate value for \sqrt{c} .

Suitable on GPU?

- Finding roots is relevant for scientific computing.
- Data parallelism: compute for many different c 's.

Application: complex root finder for polynomials in one variable.

Introduction to CUDA

1 Our first GPU Program

- running Newton's method in complex arithmetic
- examining the CUDA Compute Capability

2 CUDA Program Structure

- steps to write code for the GPU
- code to compute complex roots
- the kernel function and main program
- a scalable programming model

CUDA Compute Capability

The compute capability of an NVIDIA GPU

- is represented by a version number in the format x.y,
- identifies the features supported by the hardware.

What does it mean for the programmer? Some examples:

1.3 : double-precision floating-point operations

2.0 : synchronizing threads

3.5 : dynamic parallelism

5.3 : half-precision floating-point operations

6.0 : atomic addition operation on 64-bit floats

The compute capability is not the same as the CUDA version.

checking the card with deviceQuery on kepler

```
$ /usr/local/cuda/samples/1_Uutilities/deviceQuery/deviceQuery
/usr/local/cuda/samples/1_Uutilities/deviceQuery/deviceQuery Starting...
```

```
CUDA Device Query (Runtime API) version (CUDART static linking)
```

Detected 3 CUDA Capable device(s)

```

Device 0: "Tesla K20c"
CUDA Driver Version / Runtime Version      6.0 / 5.5
CUDA Capability Major/Minor version number: 3.5
Total amount of global memory:              4800 MBytes (5032706048 bytes)
(13) Multiprocessors, (192) CUDA Cores/MP: 2496 CUDA Cores
GPU Clock rate:                            706 MHz (0.71 GHz)
Memory Clock rate:                          2600 Mhz
Memory Bus Width:                           320-bit
L2 Cache Size:                             1310720 bytes
Maximum Texture Dimension Size (x,y,z)      1D=(65536), 2D=(65536, 65536), 3D=(4096, 4096, 4096)
Maximum Layered 1D Texture Size, (num) layers 1D=(16384), 2048 layers
Maximum Layered 2D Texture Size, (num) layers 2D=(16384, 16384), 2048 layers
Total amount of constant memory:            65536 bytes
Total amount of shared memory per block:    49152 bytes
Total number of registers available per block: 65536
Warp size:                                  32
Maximum number of threads per multiprocessor: 2048
Maximum number of threads per block:        1024
Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
Max dimension size of a grid size    (x,y,z): (2147483647, 65535, 65535)
Maximum memory pitch:                      2147483647 bytes
Texture alignment:                          512 bytes
Concurrent copy and kernel execution:       Yes with 2 copy engine(s)
Run time limit on kernels:                  No
Integrated GPU sharing Host Memory:         No
Support host page-locked memory mapping:   Yes
Alignment requirement for Surfaces:        Yes
Device has ECC support:                     Enabled
Device supports Unified Addressing (UVA):   Yes
Device PCI Bus ID / PCI location ID:       4 / 0
Compute Mode:
    < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >

```

checking the card with deviceQuery on pascal

```
$ /usr/local/cuda/samples/1_Utilities/deviceQuery/deviceQuery
/usr/local/cuda/samples/1_Utilities/deviceQuery/deviceQuery Starting...
```

```
CUDA Device Query (Runtime API) version (CUDART static linking)
```

Detected 2 CUDA Capable device(s)

```

Device 0: "Tesla P100-PCIE-16GB"
CUDA Driver Version / Runtime Version      8.0 / 8.0
CUDA Capability Major/Minor version number: 6.0
Total amount of global memory:              16276 MBytes (17066885120 bytes)
(56) Multiprocessors, ( 64) CUDA Cores/MP: 3584 CUDA Cores
GPU Max Clock rate:                        405 MHz (0.41 GHz)
Memory Clock rate:                         715 Mhz
Memory Bus Width:                          4096-bit
L2 Cache Size:                             4194304 bytes
Maximum Texture Dimension Size (x,y,z)      1D=(131072), 2D=(131072, 65536), 3D=(16384, 16384, 16384)
Maximum Layered 1D Texture Size, (num) layers 1D=(32768), 2048 layers
Maximum Layered 2D Texture Size, (num) layers 2D=(32768, 32768), 2048 layers
Total amount of constant memory:             65536 bytes
Total amount of shared memory per block:     49152 bytes
Total number of registers available per block: 65536
Warp size:                                   32
Maximum number of threads per multiprocessor: 2048
Maximum number of threads per block:         1024
Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
Max dimension size of a grid size    (x,y,z): (2147483647, 65535, 65535)
Maximum memory pitch:                       2147483647 bytes
Texture alignment:                           512 bytes
Concurrent copy and kernel execution:        Yes with 2 copy engine(s)
Run time limit on kernels:                   No
Integrated GPU sharing Host Memory:          No
Support host page-locked memory mapping:     Yes
Alignment requirement for Surfaces:          Yes
Device has ECC support:                      Enabled
Device supports Unified Addressing (UVA):    Yes
Device PCI Domain ID / Bus ID / location ID: 0 / 2 / 0
Compute Mode:
    < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >

```

running bandwidthTest on kepler

```
$ /usr/local/cuda/samples/1_Uutilities/bandwidthTest/bandwidthTest
```

```
[CUDA Bandwidth Test] - Starting...
```

```
Running on...
```

```
Device 0: Tesla K20c
```

```
Quick Mode
```

```
Host to Device Bandwidth, 1 Device(s)
```

```
PINNED Memory Transfers
```

```
Transfer Size (Bytes) Bandwidth(MB/s)
```

```
33554432 5819.5
```

```
Device to Host Bandwidth, 1 Device(s)
```

```
PINNED Memory Transfers
```

```
Transfer Size (Bytes) Bandwidth(MB/s)
```

```
33554432 6415.8
```

```
Device to Device Bandwidth, 1 Device(s)
```

```
PINNED Memory Transfers
```

```
Transfer Size (Bytes) Bandwidth(MB/s)
```

```
33554432 143248.0
```

```
Result = PASS
```


running bandwidthTest on pascal

```
$ /usr/local/cuda/samples/1_Uutilities/bandwidthTest/bandwidthTest  
[CUDA Bandwidth Test] - Starting...
```

```
Running on...
```

```
Device 0: Tesla P100-PCIE-16GB
```

```
Quick Mode
```

```
Host to Device Bandwidth, 1 Device(s)
```

```
PINNED Memory Transfers
```

```
Transfer Size (Bytes) Bandwidth(MB/s)  
33554432 11530.1
```

```
Device to Host Bandwidth, 1 Device(s)
```

```
PINNED Memory Transfers
```

```
Transfer Size (Bytes) Bandwidth(MB/s)  
33554432 12848.3
```

```
Device to Device Bandwidth, 1 Device(s)
```

```
PINNED Memory Transfers
```

```
Transfer Size (Bytes) Bandwidth(MB/s)  
33554432 444598.8
```

```
Result = PASS
```

```
$
```

Introduction to CUDA

1 Our first GPU Program

- running Newton's method in complex arithmetic
- examining the CUDA Compute Capability

2 CUDA Program Structure

- steps to write code for the GPU
- code to compute complex roots
- the kernel function and main program
- a scalable programming model

steps to write code for the GPU

Five steps to get GPU code running:

- 1 C and C++ functions are labeled with CUDA keywords `__device__`, `__global__`, or `__host__`.
- 2 Determine the data for each thread to work on.
- 3 Transferring data from/to host (CPU) to/from the device (GPU).
- 4 Statements to launch data-parallel functions, called *kernels*.
- 5 Compilation with `nvcc`.

step 1: CUDA extensions to functions

Three keywords before a function declaration:

`__host__` : The function will run on the host (CPU).

`__device__` : The function will run on the device (GPU).

`__global__` : The function is called from the host but runs on the device. This function is called a *kernel*.


CUDA extensions to C function declarations:

	executed on	callable from
<code>__device__ double D()</code>	device	device
<code>__global__ void K()</code>	device	host
<code>__host__ int H()</code>	host	host

step 2: data for each thread

The grid consists of N blocks, with `blockIdx.x` $\in \{0, N - 1\}$.

Within each block, `threadIdx.x` $\in \{0, \text{blockDim.x} - 1\}$.



```
int threadId = blockIdx.x *  
    blockDim.x + threadIdx.x  
...  
float x = input[threadId]  
float y = f(x)  
output[threadId] = y  
...
```

step 3: allocating and transferring data

```
cudaDoubleComplex *xhost = new cudaDoubleComplex[n];

// we copy n complex numbers to the device
size_t s = n*sizeof(cudaDoubleComplex);
cudaDoubleComplex *xdevice;
cudaMalloc((void**)&xdevice,s);

cudaMemcpy(xdevice,xhost,s,cudaMemcpyHostToDevice);

// allocate memory for the result
cudaDoubleComplex *ydevice;
cudaMalloc((void**)&ydevice,s);

// copy results from device to host
cudaDoubleComplex *yhost = new cudaDoubleComplex[n];

cudaMemcpy(yhost,ydevice,s,cudaMemcpyDeviceToHost);
```

step 4: launching the kernel

The kernel is declared as

```
__global__ void squareRoot
( int n, cudaDoubleComplex *x, cudaDoubleComplex *y )
// Applies Newton's method to compute the square root
// of the n numbers in x and places the results in y.
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    ...
}
```

For frequency f , dimension n , and block size w , we do:

```
// invoke the kernel with n/w blocks per grid
// and w threads per block
for(int i=0; i<f; i++)
    squareRoot<<<n/w,w>>>(n,xdevice,ydevice);
```

step 5: compiling with `nvcc`

Then if the `makefile` contains

```
runCudaComplexSqrt:  
    nvcc -o /tmp/run_cmpsqr -arch=sm_13 \  
        runCudaComplexSqrt.cu
```

typing `make runCudaComplexSqrt` does

```
nvcc -o /tmp/run_cmpsqr -arch=sm_13 runCudaComplexSqrt.cu
```

The `-arch=sm_13` is needed for double arithmetic, for the K20C.

The option `-arch=sm_13` is no longer recognized on the new P100.

Introduction to CUDA

1 Our first GPU Program

- running Newton's method in complex arithmetic
- examining the CUDA Compute Capability

2 CUDA Program Structure

- steps to write code for the GPU
- **code to compute complex roots**
- the kernel function and main program
- a scalable programming model

defining complex numbers

```
#ifndef __CUDADOUBLECOMPLEX_CU__
#define __CUDADOUBLECOMPLEX_CU__

#include <cmath>
#include <cstdlib>
#include <iomanip>
#include <vector_types.h>
#include <math_functions.h>

typedef double2 cudaDoubleComplex;
```

We use the `double2` of `vector_types.h` to define complex numbers because `double2` is a native CUDA type allowing for coalesced memory access.

random complex numbers

```
__host__ cudaDoubleComplex randomDoubleComplex()  
// Returns a complex number on the unit circle  
// with angle uniformly generated in  $[0, 2\pi]$ .  
{  
    cudaDoubleComplex result;  
    int r = rand();  
    double u = double(r)/RAND_MAX;  
    double angle = 2.0*M_PI*u;  
    result.x = cos(angle);  
    result.y = sin(angle);  
    return result;  
}
```

calling sqrt of math_functions.h

```
__device__ double radius ( const cudaDoubleComplex c )  
// Returns the radius of the complex number.  
{  
    double result;  
    result = c.x*c.x + c.y*c.y;  
    return sqrt(result);  
}
```

overloading for output

```
__host__ std::ostream& operator<<
( std::ostream& os, const cudaDoubleComplex& c)
// Writes real and imaginary parts of c,
// in scientific notation with precision 16.
{
    os << std::scientific << std::setprecision(16)
        << c.x << "  " << c.y;
    return os;
}
```

defining complex addition

```
__device__ cudaDoubleComplex operator+
( const cudaDoubleComplex a, const cudaDoubleComplex b )
// Returns the sum of a and b.
{
    cudaDoubleComplex result;
    result.x = a.x + b.x;
    result.y = a.y + b.y;
    return result;
}
```

The rest of the arithmetical operations are defined in a similar manner.

All definitions related to complex numbers are stored in the file `cudaDoubleComplex.cu`.

Introduction to CUDA

1 Our first GPU Program

- running Newton's method in complex arithmetic
- examining the CUDA Compute Capability

2 CUDA Program Structure

- steps to write code for the GPU
- code to compute complex roots
- **the kernel function and main program**
- a scalable programming model

the kernel function

```
#include "cudaDoubleComplex.cu"

__global__ void squareRoot
( int n, cudaDoubleComplex *x, cudaDoubleComplex *y )
// Applies Newton's method to compute the square root
// of the n numbers in x and places the results in y.
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    cudaDoubleComplex inc;
    cudaDoubleComplex c = x[i];
    cudaDoubleComplex r = c;
    for(int j=0; j<5; j++)
    {
        inc = r + r;
        inc = (r*r - c)/inc;
        r = r - inc;
    }
    y[i] = r;
}
```


the main function — command line arguments

```
int main ( int argc, char*argv[] )
{
    if(argc < 5)
    {
        cout << "call with 4 arguments : " << endl;
        cout << "dimension, block size, frequency, and check (0 or 1)"
            << endl;
    }
    else
    {
        int n = atoi(argv[1]); // dimension
        int w = atoi(argv[2]); // block size
        int f = atoi(argv[3]); // frequency
        int t = atoi(argv[4]); // test or not
        // we generate n random complex numbers on the host
        cudaDoubleComplex *xhost = new cudaDoubleComplex[n];
        for(int i=0; i<n; i++) xhost[i] = randomDoubleComplex();
    }
}
```

The main program generates n random complex numbers with radius 1.

transferring data and launching the kernel

```
// copy the n random complex numbers to the device
size_t s = n*sizeof(cudaDoubleComplex);
cudaDoubleComplex *xdevice;
cudaMalloc((void**)&xdevice,s);
cudaMemcpy(xdevice,xhost,s,cudaMemcpyHostToDevice);
// allocate memory for the result
cudaDoubleComplex *ydevice;
cudaMalloc((void**)&ydevice,s);
// invoke the kernel with n/w blocks per grid
// and w threads per block
for(int i=0; i<f; i++)
    squareRoot<<<n/w,w>>>(n,xdevice,ydevice);
// copy results from device to host
cudaDoubleComplex *yhost = new cudaDoubleComplex[n];
cudaMemcpy(yhost,ydevice,s,cudaMemcpyDeviceToHost);
```

testing one random number

```
if(t == 1) // test the result
{
    int k = rand() % n;
    cout << "testing number " << k << endl;
    cout << "          x = " << xhost[k] << endl;
    cout << "    sqrt(x) = " << yhost[k] << endl;
    cudaDoubleComplex z = Square(yhost[k]);
    cout << "sqrt(x)^2 = " << z << endl;
}
}
return 0;
}
```

Introduction to CUDA

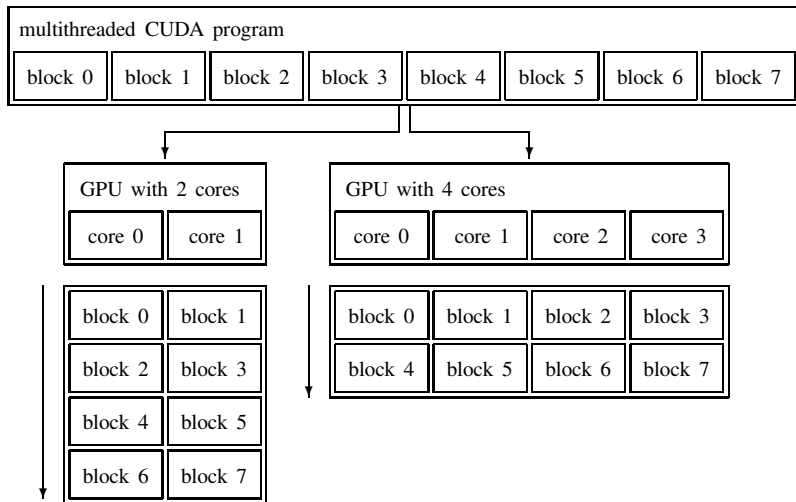
1 Our first GPU Program

- running Newton's method in complex arithmetic
- examining the CUDA Compute Capability

2 CUDA Program Structure

- steps to write code for the GPU
- code to compute complex roots
- the kernel function and main program
- a scalable programming model

a scalable programming model



running the code on kepler

A test on the correctness:

```
$ /tmp/run_cmpsqrt 1 1 1 1
testing number 0
      x = 5.3682227446949737e-01  -8.4369535119816541e-01
      sqrt(x) = 8.7659063264145631e-01  -4.8123680528950746e-01
      sqrt(x)^2 = 5.3682227446949726e-01  -8.4369535119816530e-01
```

On 64,000 numbers, 32 threads in a block, doing it 10,000 times:

```
$ time /tmp/run_cmpsqrt 64000 32 10000 1
testing number 50325
      x = 7.9510606509728776e-01  -6.0647039931517477e-01
      sqrt(x) = 9.4739275517002119e-01  -3.2007337822967424e-01
      sqrt(x)^2 = 7.9510606509728765e-01  -6.0647039931517477e-01

real    0m1.618s
user    0m0.526s
sys     0m0.841s
$
```

changing #threads in a block

```
$ time /tmp/run_cmpsqrtrt 128000 32 100000 0
```

```
real    0m17.345s
user    0m9.829s
sys     0m7.303s
```

```
$ time /tmp/run_cmpsqrtrt 128000 64 100000 0
```

```
real    0m10.502s
user    0m5.711s
sys     0m4.497s
```

```
$ time /tmp/run_cmpsqrtrt 128000 128 100000 0
```

```
real    0m9.295s
user    0m5.231s
sys     0m3.865s
```

running the code on pascal

```
$ time /tmp/run_cmpsqr 128000 32 100000 0
```

```
real    0m2.516s
user    0m1.250s
sys     0m1.236s
```

```
$ time /tmp/run_cmpsqr 128000 64 100000 0
```

```
real    0m2.521s
user    0m1.245s
sys     0m1.234s
```

```
$ time /tmp/run_cmpsqr 128000 128 100000 0
```

```
real    0m2.496s
user    0m1.288s
sys     0m1.195s
```


summary and references

In five steps we wrote our first complete CUDA program.

We started chapter 3 of the textbook by Kirk & Hwu, covering more of the CUDA Programming Guide.

Available in `/usr/local/cuda/doc` are

- CUDA C Best Practices Guide
- CUDA Programming Guide

Also available online at nvidia.com.

Many examples of CUDA applications are available in `/usr/local/cuda/samples`.

exercises

- 1 Instead of 5 Newton iterations in `runCudaComplexSqrt.cu` use k iterations where k is entered by the user at the command line. What is the influence of k on the timings?
- 2 Modify the kernel for the complex square root so it takes on input an array of complex coefficients of a polynomial of degree d . Then the root finder applies Newton's method, starting at random points. Test the correctness and experiment to find the rate of success, i.e.: for polynomials of degree d how many random trials are needed to obtain $d/2$ roots of the polynomial?