

# Quad Doubles on a GPU

## 1 Floating-Point Arithmetic

- floating-point numbers
- quad double arithmetic
- quad doubles for use in CUDA programs

## 2 Quad Double Square Roots

- quad double arithmetic on a GPU
- a kernel using `gqd_real`
- performance considerations

MCS 572 Lecture 37  
Introduction to Supercomputing  
Jan Verschelde, 16 November 2016

# Quad Doubles on a GPU

## 1 Floating-Point Arithmetic

- floating-point numbers
- quad double arithmetic
- quad doubles for use in CUDA programs

## 2 Quad Double Square Roots

- quad double arithmetic on a GPU
- a kernel using `gqd_real`
- performance considerations

# floating-point numbers

A floating-point number consists of a sign bit, exponent, and a fraction (also known as the mantissa).

Almost all microprocessors follow the IEEE 754 standard.

GPU hardware supports 32-bit (single float)  
and for compute capability  $\geq 1.3$  also double floats.

Numerical analysis studies algorithms for continuous problems,  
investigating

- problems for their sensitivity to errors in the input; and
- algorithms for their propagation of roundoff errors.

# parallel numerical algorithms

The floating-point addition is *not* associative!

Parallel algorithms compute and accumulate the results in an order that is different from their sequential versions.

Example: Adding a sequence of numbers is more accurate if the numbers are sorted in increasing order.

Instead of speedup, we can ask questions about quality up:

- If we can afford to keep the total running time constant, does a faster computer give us more accurate results?
- How many more processors do we need to guarantee a result?

# Quad Doubles on a GPU

## 1 Floating-Point Arithmetic

- floating-point numbers
- **quad double arithmetic**
- quad doubles for use in CUDA programs

## 2 Quad Double Square Roots

- quad double arithmetic on a GPU
- a kernel using `gqd_real`
- performance considerations

# quadruple precision

A quad double is an unevaluated sum of 4 doubles, improves working precision from  $2.2 \times 10^{-16}$  to  $2.4 \times 10^{-63}$ .

Y. Hida, X.S. Li, and D.H. Bailey: **Algorithms for quad-double precision floating point arithmetic**. In *15th IEEE Symposium on Computer Arithmetic* pages 155–162. IEEE, 2001. Software at <http://crd.lbl.gov/~dhbailey/mpdist>.

A quad double builds on `double double`, some features:

- The least significant part of a `double double` can be interpreted as a compensation for the roundoff error.
- Predictable overhead: working with `double double` is of the same cost as working with complex numbers.

# Newton's method for $\sqrt{x}$

```
#include <iostream>
#include <iomanip>
#include <qd/qd_real.h>
using namespace std;

qd_real newton ( qd_real x )
{
    qd_real y = x;
    for(int i=0; i<10; i++)
        y -= (y*y - x) / (2.0*y);
    return y;
}
```

# the main program

```
int main ( int argc, char *argv[] )
{
    cout << "give x : ";
    qd_real x; cin >> x;
    cout << setprecision(64);
    cout << "          x : " << x << endl;

    qd_real y = newton(x);
    cout << "  sqrt(x) : " << y << endl;

    qd_real z = y*y;
    cout << "sqrt(x)^2 : " << z << endl;

    return 0;
}
```



## the makefile

If the program is on file `newton4sqrt.cpp`  
and the makefile contains

```
QD_ROOT=/usr/local/qd-2.3.17
```

```
QD_LIB=/usr/local/lib
```

```
newton4sqrt:
```

```
    g++ -I$(QD_ROOT)/include newton4sqrt.cpp \  
        $(QD_LIB)/libqd.a -o /tmp/newton4sqrt
```

then we can create the executable as

```
$ make newton4sqrt
```

```
g++ -I/usr/local/qd-2.3.17/include newton4sqrt.cpp \  
    /usr/local/lib/libqd.a -o /tmp/newton4sqrt
```

```
$
```

# Quad Doubles on a GPU

## 1 Floating-Point Arithmetic

- floating-point numbers
- quad double arithmetic
- quad doubles for use in CUDA programs

## 2 Quad Double Square Roots

- quad double arithmetic on a GPU
- a kernel using `gqd_real`
- performance considerations

# extended precision on the GPU

Large problems often need extra precision.

The QD library has been ported to the GPU.

Mian Lu, Bingsheng He, and Qiong Luo: **Supporting extended precision on graphics processors**. In A. Ailamaki and P.A. Boncz, editors, *Proceedings of the Sixth International Workshop on Data Management on New Hardware (DaMoN 2010)*, June 7, 2010, Indianapolis, Indiana, pages 19–26, 2010.

Software at <http://code.google.com/p/gpuprec/>, and at <https://github.com/lumianph/gpuprec/tree/master/gqd>.

Installed on `kepler` and `pascal` in `/usr/local/gqd_1_2`.

For graphics cards of compute capability  $< 1.3$ , one could use the freely available Cg software of Andrew Thall to achieve double precision using float-float arithmetic.

## gqd\_reals are of double4 type

```
#include "gqd_type.h"
#include "vector_types.h"
#include <qd/qd_real.h>

void qd2gqd ( qd_real *a, gqd_real *b )
{
    b->x = a->x[0];
    b->y = a->x[1];
    b->z = a->x[2];
    b->w = a->x[3];
}

void gqd2qd ( gqd_real *a, qd_real *b )
{
    b->x[0] = a->x;
    b->x[1] = a->y;
    b->x[2] = a->z;
    b->x[3] = a->w;
}
```

## a first kernel

```
#include "gqd.cu"
```

```
__global__ void testdiv2 ( gqd_real *x, gqd_real *y )  
{  
    *y = *x/2.0;  
}
```

```
int divide_by_two ( gqd_real *x, gqd_real *y )  
{  
    gqd_real *xdevice;  
    size_t s = sizeof(gqd_real);  
    cudaMalloc((void**)&xdevice,s);  
    cudaMemcpy(xdevice,x,s,cudaMemcpyHostToDevice);  
    gqd_real *ydevice;  
    cudaMalloc((void**)&ydevice,s);  
    testdiv2<<<1,1>>>(xdevice,ydevice);  
    cudaMemcpy(y,ydevice,s,cudaMemcpyDeviceToHost);  
    return 0;  
}
```

## testing the first kernel

```
#include <iostream>
#include <iomanip>
#include "gqd_type.h"
#include "first_gqd_kernel.h"
#include "gqd_qd_util.h"
#include <qd/qd_real.h>
using namespace std;

int main ( int argc, char *argv[] )
{
    qd_real qd_x = qd_real::_pi;
    gqd_real x;
    qd2gqd(&qd_x, &x);
    gqd_real y;

    cout << " x : " << setprecision(64) << qd_x << endl;
```

## test program continued

```
int fail = divide_by_two(&x,&y);

qd_real qd_y;
gqd2qd(&y,&qd_y);

if(fail == 0) cout << " y : " << qd_y << endl;

cout << "2y : " << 2.0*qd_y << endl;

return 0;
}
```

# the makefile

```
QD_ROOT=/usr/local/qd-2.3.17
```

```
QD_LIB=/usr/local/lib
```

```
GQD_HOME=/usr/local/gqd_1_2
```

```
SDK_HOME=/usr/local/cuda/sdk
```

```
test_pi2_gqd_kernel:
```

```
    @-echo ">>> compiling kernel ..."
```

```
    nvcc -I$(GQD_HOME)/inc -I$(SDK_HOME)/C/common/inc \
        -c first_gqd_kernel.cu
```

```
    @-echo ">>> compiling utilities ..."
```

```
    g++ -I/usr/local/cuda/include -I$(GQD_HOME)/inc \
        -I$(QD_ROOT)/include -c gqd_qd_util.cpp
```

```
    @-echo ">>> compiling test program ..."
```

```
    g++ test_pi2_gqd_kernel.cpp -c \
        -I/usr/local/cuda/include -I$(GQD_HOME)/inc \
        -I$(QD_ROOT)/include
```

```
    @-echo ">>> linking ..."
```

```
    g++ -I$(GQD_HOME)/inc -I$(QD_ROOT)/include \
        first_gqd_kernel.o test_pi2_gqd_kernel.o gqd_qd_util.o \
        $(QD_LIB)/libqd.a \
        -o /tmp/test_pi2_gqd_kernel \
        -lcuda -lcutil_x86_64 -lcudart \
        -L/usr/local/cuda/lib64 -L$(SDK_HOME)/C/lib
```



# compiling and running

```
$ make test_pi2_gqd_kernel
>>> compiling kernel ...
nvcc -I/usr/local/gqd_1_2/inc -I/usr/local/cuda/sdk/C/common/inc \
      -c first_gqd_kernel.cu
>>> compiling utilities ...
g++ -I/usr/local/cuda/include -I/usr/local/gqd_1_2/inc \
      -I/usr/local/qd-2.3.17/include -c gqd_qd_util.cpp
>>> compiling test program ...
g++ test_pi2_gqd_kernel.cpp -c \
      -I/usr/local/cuda/include -I/usr/local/gqd_1_2/inc \
      -I/usr/local/qd-2.3.17/include
>>> linking ...
g++ -I/usr/local/gqd_1_2/inc -I/usr/local/qd-2.3.17/include \
      first_gqd_kernel.o test_pi2_gqd_kernel.o gqd_qd_util.o \
      /usr/local/lib/libqd.a \
      -o /tmp/test_pi2_gqd_kernel \
      -lcuda -lcutil_x86_64 -lcudart \
      -L/usr/local/cuda/lib64 -L/usr/local/cuda/sdk/C/lib
$ /tmp/test_pi2_gqd_kernel
  x : 3.1415926535897932384626433832795028841971693993751058209749445923e+00
  y : 1.5707963267948966192313216916397514420985846996875529104874722961e+00
2y : 3.1415926535897932384626433832795028841971693993751058209749445923e+00
$
```

# Quad Doubles on a GPU

## 1 Floating-Point Arithmetic

- floating-point numbers
- quad double arithmetic
- quad doubles for use in CUDA programs

## 2 Quad Double Square Roots

- quad double arithmetic on a GPU
- a kernel using `gqd_real`
- performance considerations

# quad double arithmetic on a GPU

Recall our first CUDA program to take the square root of complex numbers stored in a `double2` array.

In using quad doubles on a GPU, we have 3 stages:

- 1 The kernel in a file with extension `cu` is compiled with `nvcc -c` into an object file.
- 2 The application code is compiled with `g++ -c`.
- 3 The linker `g++` takes `.o` files and libraries on input to make an executable file.

Working without a makefile now becomes very tedious.

# the makefile

```
QD_ROOT=/usr/local/qd-2.3.17
```

```
QD_LIB=/usr/local/lib
```

```
GQD_HOME=/usr/local/gqd_1_2
```

```
SDK_HOME=/usr/local/cuda/sdk
```

```
sqrt_gqd_kernel:
```

```
    @-echo ">>> compiling kernel ..."
```

```
    nvcc -I$(GQD_HOME)/inc -I$(SDK_HOME)/C/common/inc \
        -c sqrt_gqd_kernel.cu
```

```
    @-echo ">>> compiling utilities ..."
```

```
    g++ -I/usr/local/cuda/include -I$(GQD_HOME)/inc \
        -I$(QD_ROOT)/include -c gqd_qd_util.cpp
```

```
    @-echo ">>> compiling test program ..."
```

```
    g++ run_sqrt_gqd_kernel.cpp -c \
        -I/usr/local/cuda/include -I$(GQD_HOME)/inc \
        -I$(QD_ROOT)/include
```

```
    @-echo ">>> linking ..."
```

```
    g++ -I$(GQD_HOME)/inc -I$(QD_ROOT)/include \
        sqrt_gqd_kernel.o run_sqrt_gqd_kernel.o gqd_qd_util.o \
        $(QD_LIB)/libqd.a \
        -o /tmp/run_sqrt_gqd_kernel \
        -lcuda -lcutil_x86_64 -lcudart \
        -L/usr/local/cuda/lib64 -L$(SDK_HOME)/C/lib
```

# running make

```
$ make sqrt_gqd_kernel
>>> compiling kernel ...
nvcc -I/usr/local/gqd_1_2/inc -I/usr/local/cuda/sdk/C/common/inc \
      -c sqrt_gqd_kernel.cu
>>> compiling utilities ...
g++ -I/usr/local/cuda/include -I/usr/local/gqd_1_2/inc \
      -I/usr/local/qd-2.3.17/include -c gqd_qd_util.cpp
>>> compiling test program ...
g++ run_sqrt_gqd_kernel.cpp -c \
      -I/usr/local/cuda/include -I/usr/local/gqd_1_2/inc \
      -I/usr/local/qd-2.3.17/include
>>> linking ...
g++ -I/usr/local/gqd_1_2/inc -I/usr/local/qd-2.3.17/include \
      sqrt_gqd_kernel.o run_sqrt_gqd_kernel.o gqd_qd_util.o \
      /usr/local/lib/libqgd.a \
      -o /tmp/run_sqrt_gqd_kernel \
      -lcuda -lcutil_x86_64 -lcudart \
      -L/usr/local/cuda/lib64 -L/usr/local/cuda/sdk/C/lib

$
```

# Quad Doubles on a GPU

## 1 Floating-Point Arithmetic

- floating-point numbers
- quad double arithmetic
- quad doubles for use in CUDA programs

## 2 Quad Double Square Roots

- quad double arithmetic on a GPU
- **a kernel using `gqd_real`**
- performance considerations

## a kernel using gqd\_real

```
#include "gqd.cu"

__global__ void sqrtNewton ( gqd_real *x, gqd_real *y )
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    gqd_real c = x[i];
    gqd_real r = c;
    for(int j=0; j<10; j++)
        r = r - (r*r - c)/(2.0*r);
    y[i] = r;
}
```

## the file `sqrt_gqd_kernel.cu` continued

```
int sqrt_by_Newton ( int n, gqd_real *x, gqd_real *y )
{
    gqd_real *xdevice;
    size_t s = n*sizeof(gqd_real);
    cudaMalloc((void**)&xdevice,s);
    cudaMemcpy(xdevice,x,s,cudaMemcpyHostToDevice);

    gqd_real *ydevice;
    cudaMalloc((void**)&ydevice,s);

    sqrtNewton<<<n/32,32>>>(xdevice,ydevice);

    cudaMemcpy(y,ydevice,s,cudaMemcpyDeviceToHost);

    return 0;
}
```



# the main program

```
#include <iostream>
#include <iomanip>
#include <cstdlib>
#include "gqd_type.h"
#include "sqrt_gqd_kernel.h"
#include "gqd_qd_util.h"
#include <qd/qd_real.h>
using namespace std;

int main ( int argc, char *argv[] )
{
    const int n = 256;
    gqd_real *x = (gqd_real*)calloc(n,sizeof(gqd_real));
    gqd_real *y = (gqd_real*)calloc(n,sizeof(gqd_real));
```

## run\_sqrt\_gqd\_kernel.cpp continued

```
for(int i = 0; i<n; i++)
{
    x[i].x = (double) (i+2);
    x[i].y = 0.0; x[i].z = 0.0; x[i].w = 0.0;
}
int fail = sqrt_by_Newton(n,x,y);
if(fail == 0)
{
    const int k = 24;
    qd_real qd_x;
    gqd2qd(&x[k],&qd_x);
    qd_real qd_y;
    gqd2qd(&y[k],&qd_y);
    cout << "      x      : " << setprecision(64) << qd_x << endl;
    cout << "sqrt(x)    : " << setprecision(64) << qd_y << endl;
    cout << "sqrt(x)^2 : " << setprecision(64) << qd_y*qd_y
        << endl;
}
return 0;
}
```

# Quad Doubles on a GPU

## 1 Floating-Point Arithmetic

- floating-point numbers
- quad double arithmetic
- quad doubles for use in CUDA programs

## 2 Quad Double Square Roots

- quad double arithmetic on a GPU
- a kernel using `gqd_real`
- **performance considerations**

# sequential and interval memory layout

Consider four quad doubles  $a$ ,  $b$ ,  $c$ , and  $d$ .

Stored in a sequential memory layout:

$a_0$	$a_1$	$a_2$	$a_3$	$b_0$	$b_1$	$b_2$	$b_3$	$c_0$	$c_1$	$c_2$	$c_3$	$d_0$	$d_1$	$d_2$	$d_3$
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15

Stored in an interval memory layout:

$a_0$	$b_0$	$c_0$	$d_0$	$a_1$	$b_1$	$c_1$	$d_1$	$a_2$	$b_2$	$c_2$	$d_2$	$a_3$	$b_3$	$c_3$	$d_3$
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15

The implementation with an interval memory layout is reported to be three times faster over the sequential memory layout.

# Bibliography

- A. Thall. **Extended-Precision Floating-Point Numbers for GPU Computation.** Software available at [andrewthall.org](http://andrewthall.org).
- Y. Hida, X.S. Li, and D.H. Bailey. **Algorithms for quad-double precision floating point arithmetic.** In *15th IEEE Symposium on Computer Arithmetic*, pages 155–162. IEEE, 2001.  
Software at <http://crd.lbl.gov/~dhbailey/mpdist>.
- M. Lu, B. He, and Q. Luo. **Supporting extended precision on graphics processors.** In A. Ailamaki and P.A. Boncz, editors, *Proceedings of the Sixth International Workshop on Data Management on New Hardware (DaMoN 2010), June 7, 2010, Indianapolis, Indiana*, pages 19–26, 2010.  
Software at <http://code.google.com/p/gpuprec/> and at <https://github.com/lumianph/gpuprec/tree/master/gqd>.

# summary and exercises

Chapter 7 in the book of Kirk & Hwu provides some background. The application of quad double arithmetic is an illustration of combined usage of `nvcc` and `g++` to compile and link several libraries.

Some exercises:

- 1 Compare the performance of the CUDA program for Newton's method for square root with quad doubles to the code of lecture 29.
- 2 Extend the code so it works for complex quad double arithmetic.
- 3 Use quad doubles to implement the second parallel sum algorithm of lecture 33. Could the parallel implementation with quad doubles run as fast as sequential code with doubles?
- 4 Consider the program to approximate  $\pi$  of lecture 13. Write a version for the GPU and compare the performance with the multicore version of lecture 13.