

# Google Interview Questions

Add Question (/add)

## Page:

1 2 (<https://www.careercup.com/page?pid=google-interview-questions&topic=system-design-interview-questions&n=2>)

**Filter:** Google (None) System Design Go Clear

**Sort By:**

Date | [Number of Comments \(?pid=google-interview-questions&topic=system-design-interview-questions&n=1&sort=comments\)](#) | [Most Recent Comment \(?pid=google-interview-questions&topic=system-design-interview-questions&n=1&sort=recentcomments\)](#) | [Votes \(?pid=google-interview-questions&topic=system-design-interview-questions&n=1&sort=votes\)](#)

**CareerCup** is the world's biggest and best source for software engineering interview preparation. See all our resources (</resources>).

Open Chat in New Window  
(<http://tlk.io/careercup>)



▲ 0 ▼  
of 0 votes

Given a large file with sentences and query string, design a system (Class, data structs, functions, etc) and algorithm to return the smallest window (start and end offsets) in the input file where the query words (in any order) are seen in the text file. What is the time complexity?

5  
Answers  
(/question?id=)

(/question?id=5720301565378560) - Ray  
 (/user?id=6011859920486400) May 01, 2017  
 (2017-05-01T18:22:19.726440Z) in United States | Report Duplicate  
 (/reportduplicate?id=5720301565378560) | Flag  
 (/flagquestion?id=5720301565378560)



▲ 0 ▼  
of 0 votes

How do you design a system for very large graphs(does not fit in a single machine)?

(/question?id=5662125751009280) - Matt Chad (/user?id=19691734)  
January 31, 2016 (2016-01-31T02:55:39.473540Z) in United States | Report  
Duplicate (/reportduplicate?id=5662125751009280) | Flag  
(/flagquestion?id=5662125751009280)

10  
Answers  
(/question?id=)

How Can You Help CareerCup?  
(/helpcareercup)

Find Out When Gayle / CareerCup / Cracking the Coding Interview is in

## Videos

CareerCup's interview videos give you a real-life look at technical interviews. In these unscripted videos, watch how other candidates handle tough questions and how the interviewer thinks about their performance.



# Mock Interviews

2) Our Mock Interviews will be conducted "in character" just like a real interview, and can focus on whatever topics you want. All our interviewers have worked for Microsoft, Google or Amazon, you know you'll get a true-to-life lender experience.



lendar experience.



ture

2) - tested.candidate

015

.....

d States | Report Duplicate

## Products

Coding Interview      Mock Interviews  
www.amazon.com      Resume Review (/r/  
t/0984782850      Acquisition Interview  
s?tl?ie=UTF8&      (http://www.gayle.com)



&linkCode=as2&  
=217145&  
85050tiveASIN=09<  
e PM interview  
015  
amazon.com  
1/60147810  
Port Duplicate  
pJf7ag=UTF8&  
O&linkCode=as2&  
&creative=399373&



```

#email=arnold@ph
t/1118968085
S, ti?ie=UTF8&
&link_code=as2&
&ref=399373&
#state=FL
#-state=FL Duplicate
$ (/misc)

```



for system for a building with  
ize in terms of lowest wait  
ed to power consumption.

B) - tbag (/user?id=7691663) March  
) in United States | Report Duplicate  
) | Flag

6  
Answers  
(/question?id=)

1  
Answer  
(/question?id=)

5  
Answers  
(/question?id=)

15  
Answers  
(/question?id=)

“loved it” “invaluable training”  
“very comprehensive” “excellent”  
“highly recommended”  
“best resource for interview preparation”

<http://www.twitter.com/careercup>)  
 NVIDIA (2013)  
 Gayle on Twitter  
 (/page?id=nvidia)  
 (http://www.twitter.com/gayle)  
 interview questions  
 (http://blog.careercup.com)

8/1/18, 10:14 AM  
Software Engineer in Test  
(820)

▲ 0 ▼  
of 0 votes

design a video thumb up/down system at youtube scale. how to concurrent read/write, persistence, store, update....etc.

(/question?id=5678600256749568) - rjrush (/user?id=20688670)  
January 30, 2015 (2015-01-30T11:29:47.733740Z) in United States | Report Duplicate (/reportduplicate?id=5678600256749568) | Flag  
(/flagquestion?id=5678600256749568)

▲ 0 ▼  
of 2 votes

Given n number of legacy services with user data<userid, info, date>  
Design an API to return user data in a given date range, it should collect data from each service and merge and return the data sorted by date.

(/question?id=5288170230906880) - next\_big\_gig (/user?id=134679)  
December 05, 2014 (2014-12-05T17:57:45.786910Z) in United States | Report Duplicate (/reportduplicate?id=5288170230906880) | Flag  
(/flagquestion?id=5288170230906880)

▲ 0 ▼  
of 2 votes

Need to implement something like pastebin wherein you paste some text, you are given an url. The url can be used anywhere to access the text.  
Various problems, features and design of this architecture were discussed.

(/question?id=5701782818783232) - kenadams.awesome  
(/user?id=5635124926349312) May 18, 2014  
(2014-05-18T20:38:55.220440Z) in India | Report Duplicate (/reportduplicate?id=5701782818783232) | Flag  
(/flagquestion?id=5701782818783232)

▲ -4 ▼  
of 8 votes

What would happen if you have only one server for a web cache (a web browser cache whose key is url and value is the loaded content of the webpage) but huge numbers of clients? And how would you solve it? Assume the cache is implemented with a hashmap and a linkedlist.

(/question?id=4807591515389952) - Guy  
(/user?id=5668113790009344) March 06, 2014  
(2014-03-06T17:57:53.434760Z) in United States | Report Duplicate (/reportduplicate?id=4807591515389952) | Flag  
(/flagquestion?id=4807591515389952)

▲ -5 ▼  
of 7 votes

How would you split a search query across multiple machines?

(/question?id=5377673471721472) - Guy  
(/user?id=5668113790009344) February 10, 2014  
(2014-02-10T03:54:43.356540Z) in United States | Report Duplicate (/reportduplicate?id=5377673471721472) | Flag  
(/flagquestion?id=5377673471721472)

▲ -3 ▼  
of 5 votes

How does trie handle scalability as opposed to hashtable? Assuming it is used for a dictionary. Scalability here should cover large size of input, running out of memory, or even running out of memory on multiple machines if distributed system is used.

(/question?id=6577142510911488) - Guy  
(/user?id=5668113790009344) February 06, 2014  
(2014-02-06T04:27:38.429910Z) in United States | Report Duplicate (/reportduplicate?id=6577142510911488) | Flag

6

Answers  
(/question?id=)

Software Engineer (642)  
(/page?pid=software-engineer-interview-questions)

SDE-2 (563)  
(/page?pid=sde-2-interview-questions)

7

Answers  
(/question?id=)

Developer Program Engineer (459)  
(/page?pid=developer-program-engineer-interview-questions)

Software Developer (383)  
(/page?pid=software-developer-interview-questions)

4

Answers  
(/question?id=)

Financial Software Developer (292)  
(/page?pid=financial-software-developer-interview-questions)

Java Developer (291)  
(/page?pid=java-developer-interview-questions)

Intern (267)  
(/page?pid=intern-interview-questions)

9

Answers  
(/question?id=)

**More Jobs »**  
(/categories)

## Top Topics

Algorithm (5952)  
(/page?pid=algorithm-interview-questions)

google (1083)  
(/page?pid=google-interview-questions)

Coding (970)  
(/page?pid=coding-interview-questions)

Data Structures (844)  
(/page?pid=data-structures-interview-questions)

Java (800)  
(/page?pid=java-interview-questions)

8/1/18, 10:14 AM  
C (575) (/page?pid=c-)

▲ -1 ▼  
of 5 votes

What happens during and after a query is being typed (autocomplete) in a search box whether the user is trying to go to a website or asking a question etc, and how do servers complete the request and what is the best (parallel) structure for the request to go through. DFS and how servers are located for proximity

(/question?id=5795040777469952) - Guy  
(/user?id=5668113790009344) January 30, 2014  
(2014-01-30T23:35:36.761590Z) in United States | Report Duplicate  
(/reportduplicate?id=5795040777469952) | Flag  
(/flagquestion?id=5795040777469952)

▲ -3 ▼  
of 5 votes

The setup is that we are given a series of text files which contain information regarding a code repository's commits. Each file represents a single commit and they are formatted as follows:

```
"
Commit #: XXX
Author: XXX
Reviewer(s): XXX, XXX, ...
File: XXX
File: XXX
...
Date: XX:XX:XX XX/XX/XXXX
"
```

The commit number is unique and is generated in synchronous order. There is exactly 1 unique author. There are a variable number of reviewers, delimited by commas; if there are no reviewers, that line is absent from the file. There are a variable number of edited files in the commit, each receiving its own line. The time/date is when the commit was submitted.

First design a graphical model for all of the commit data. Then describe how this model is updated when a new commit is generated. Finally, write the code segment called when a new commit is generated which edits a system that has implemented your model of the data - its input is a file name and whatever necessary data structures that are maintained by your system.

(/question?id=5105977028247552) - Guy  
(/user?id=5668113790009344) January 29, 2014  
(2014-01-29T22:35:00.003660Z) in United States | Report Duplicate  
(/reportduplicate?id=5105977028247552) | Flag  
(/flagquestion?id=5105977028247552)

▲ 1 ▼  
of 3 votes

How would you synchronize a linked list across multiple computers. If nodes are added/removed to a linked list on one computer, all others must reflect this change. Concurrency must be accounted for. (java)

(/question?id=5695609600212992) - Guy  
(/user?id=5668113790009344) January 22, 2014  
(2014-01-22T06:54:41.084400Z) in United States | Report Duplicate  
(/reportduplicate?id=5695609600212992) | Flag  
(/flagquestion?id=5695609600212992)

4 of 6  
of 3 votes

Design Short URL. (I am not sure what it even means)

(/question?id=4647544969756672) - Guy  
(/user?id=5668113790009344) January 18, 2014  
(2014-01-18T20:00:59.324120Z) in United States | Report Duplicate

3

Answers  
(/question?id=)

interview questions)  
C++ (549) (/page?pid=c-plus-plus-interview-questions)

Arrays (458)  
(/page?pid=arrays-interview-questions)

Object Oriented Design (356) (/page?pid=object-oriented-design-interview-questions)

Brain Teasers (329)  
(/page?pid=brain-teasers-interview-questions)

**More Topics »**  
(/categories)

8

Answers  
(/question?id=)

## What's Going On



**lolbnja121** said  
<http://hpprintersupportno.com>



**markjohnn09** said Epson printers are the most advanced of all other printing ...



**markjohnn09** said Epson printers are the most advanced of all other printing ...



Report a Bug or Issue  
(<https://careercup.wufoo.cc/forms/careercup-bug-report/>)

18

Answers  
(/question?id=)

8

Answers  
(/question?id=)





▲ 0 ▼  
of 2 votes

How would you model the animal kingdom (with species and their behavior) as a class system?

(/question?id=5731487106203648) - Guy  
 (/user?id=5668113790009344) January 08, 2014  
 (2014-01-08T16:42:50.086180Z) in United States | Report Duplicate  
 (/reportduplicate?id=5731487106203648) | Flag  
 (/flagquestion?id=5731487106203648)

4  
Answers  
(/question?id=5731487106203648)



▲ -3 ▼  
of 5 votes

How would you model the animal kingdom (with species and their behavior) as a class system? Java

(/question?id=5948026472366080) - Guy  
 (/user?id=5668113790009344) January 08, 2014  
 (2014-01-08T16:32:09.060690Z) in United States | Report Duplicate  
 (/reportduplicate?id=5948026472366080) | Flag  
 (/flagquestion?id=5948026472366080)

0  
Answers  
(/question?id=5948026472366080)



▲ -1 ▼  
of 3 votes

How would you model the animal kingdom (with species and their behavior) as a class system?

(/question?id=5464585372958720) - Guy  
 (/user?id=5668113790009344) January 07, 2014  
 (2014-01-07T20:50:07.907990Z) in United States | Report Duplicate  
 (/reportduplicate?id=5464585372958720) | Flag  
 (/flagquestion?id=5464585372958720)

5  
Answers  
(/question?id=5464585372958720)



▲ 0 ▼  
of 0 votes

Design a Rubik's Cube, including backend database portion.

(/question?id=6538761173204992) - asiamgenius  
 (/user?id=5733374408785920) December 27, 2013  
 (2013-12-27T21:10:11.137560Z) in United States | Report Duplicate  
 (/reportduplicate?id=6538761173204992) | Flag  
 (/flagquestion?id=6538761173204992)

7  
Answers  
(/question?id=6538761173204992)



▲ 4 ▼  
of 4 votes

design a system to return an unique ID for each request. For most of requests, the ID value should increase as time goes, the system should handle 1000 requests per second at least. timestamps alone is not valid since there might be multiple requests with same timestamps.

(/question?id=5169800024162304) - codingfunnyguy  
 (/user?id=5161939730694144) December 04, 2013  
 (2013-12-04T02:51:41.169890Z) in United States | Report Duplicate  
 (/reportduplicate?id=5169800024162304) | Flag  
 (/flagquestion?id=5169800024162304)

23  
Answers  
(/question?id=5169800024162304)



▲ 22 ▼  
of 24 votes

Given a large network of computers, each keeping log files of visited urls, find the top ten of the most visited urls. (i.e. have many large <string (url) -> int (visits)> maps, calculate implicitly <string (url) -> int (sum of visits among all distributed maps), and get the top ten in the combined map)

The result list must be exact, and the maps are too large to transmit over the network (especially sending all of them to a central server or using MapReduce directly, is not allowed)

66  
Answers  
(/question?id=5169800024162304)



▲ 0 ▼  
of 0 votes

How do you make sure an API does not leak memory?

(/question?id=16359670) - Romonov (/user?id=16008674) March 19, 2013 (2013-03-19T16:07:35.555117Z) in United States | Report Duplicate  
(/reportduplicate?id=16359670) | Flag (/flagquestion?id=16359670)

16  
Answers  
(/question?id=16359670)



▲ 0 ▼  
of 0 votes

design a server architecture for serving Google maps images

(/question?id=13148675) - GKalchev (/user?id=12932664) March 23, 2012 (2012-03-23T15:26:26.143162Z) in United States | Report Duplicate  
(/reportduplicate?id=13148675) | Flag (/flagquestion?id=13148675)

2  
Answers  
(/question?id=13148675)



▲ 0 ▼  
of 0 votes

Design a distributed system for storing a static set of (key,value) pairs and should serve user-queries(users provide a key, and system should return a the value).  
This was a system design question where I was supposed to find the number of m/cs required, networking between them, hard disks, distributing the values across them, identify bottlenecks, replication, reliability, estimation of response time, throughput etc.

For designing the system, he gave me set of resources, like configuration of the machines available, Ethernet cable B/w, etc

(/question?id=8706680) - Anonymous April 28, 2011 (2011-04-28T23:13:47.251870Z) | Report Duplicate  
(/reportduplicate?id=8706680) | Flag (/flagquestion?id=8706680)

8  
Answers  
(/question?id=8706680)



▲ -1 ▼  
of 3 votes

how you will design system for server and which will have many clients, all clients will keep of adding some words to system and server has to detect which words to accept, at the same time server will also keep on adding words. So which data structure you will use so that system will be efficient and fast.

(/question?id=7391666) - ashish January 09, 2011 (2011-01-09T05:56:07.384868Z) | Report Duplicate  
(/reportduplicate?id=7391666) | Flag (/flagquestion?id=7391666)

8  
Answers  
(/question?id=7391666)



▲ 0 ▼  
of 0 votes

Design a logging system for an application server? see to it that Logging system you define does not include a large overhead in case of large loads to server ?

(/question?id=316717) - G January 29, 2010 (2010-01-29T02:23:43.586657Z) | Report Duplicate  
(/reportduplicate?id=316717) | Flag (/flagquestion?id=316717)

1  
Answer  
(/question?id=316717)

#### Page:

1 2 (<https://www.careercup.com/page?pid=google-interview-questions&topic=system-design-interview-questions&n=2>)

#### Filter:

(None)

(None)

(None)

Go

Clear

#### Sort By:

Date | [Number of Comments \(?pid=google-interview-questions&topic=system-design-interview-questions&n=1&sort=comments\)](#) | [Most Recent Comment \(?pid=google-interview-questions&topic=system-design-interview-questions&n=1&sort=recentcomments\)](#) | [Votes \(?pid=google-interview-questions&topic=system-design-interview-questions&n=1&sort=votes\)](#)