

# Device Memories and Matrix Multiplication

## 1 Device Memories

- global, constant, and shared memories
- CUDA variable type qualifiers

## 2 Matrix Multiplication

- an application of tiling
- running `matrixMul` in the GPU Computing SDK
- the kernel of `matrixMul`

MCS 572 Lecture 32  
Introduction to Supercomputing  
Jan Verschelde, 4 November 2016

# Device Memories and Matrix Multiplication

## 1 Device Memories

- global, constant, and shared memories
- CUDA variable type qualifiers

## 2 Matrix Multiplication

- an application of tiling
- running `matrixMul` in the GPU Computing SDK
- the kernel of `matrixMul`

# data storage on a graphics card

Before we launch a kernel, we have

- to allocate memory on the device,
- to transfer data from the host to the device.

By default, memory on the device is *global memory*.

In addition to global memory, we distinguish between

- registers for storing local variables,
- shared memory for all threads in a block,
- constant memory for all blocks on a grid.

# compute to global memory access (CGMA) ratio

The importance of understanding different memories is in the calculation of the expected performance level of kernel code.

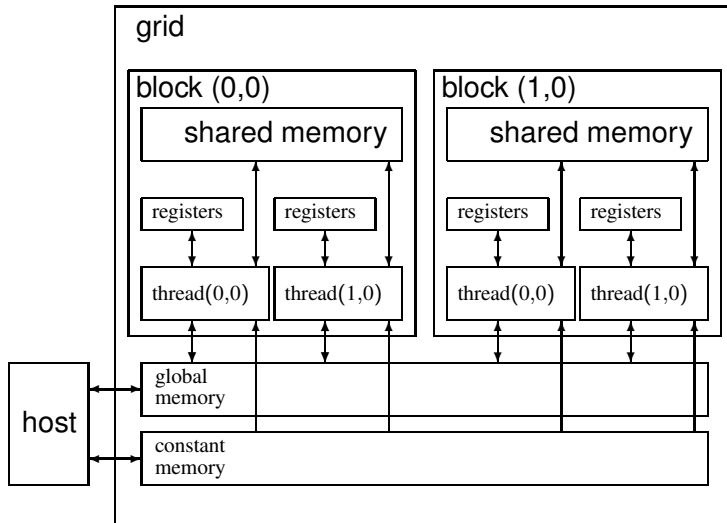
## Definition

The *Compute to Global Memory Access (CGMA) ratio* is the number of floating-point calculations performed for each access to the global memory within a region of a CUDA program.

If the CGMA ratio is 1.0, then the memory clock rate determines the upper limit for the performance.

While memory bandwidth on a GPU is superior to that of a CPU, we will miss the theoretical peak performance by a factor of ten.

# CUDA device memory types



# registers

Registers are allocated to individual threads.  
Each thread can access only its own registers.

A kernel function typically uses registers to hold frequently accessed variables that are private to each thread.

Number of 32-bit registers available per block:

- 8,192 on the GeForce 9400M,
- 32,768 on the Tesla C2050/C2070,
- 65,536 on the Tesla K20C and the P100.

A typical CUDA kernel may launch thousands of threads.

However, having too many local variables in a kernel function may prevent all blocks from running in parallel.

# shared memory

Like registers, shared memory is an on-chip memory.

Variables residing in registers and shared memory can be accessed at very high speed in a highly parallel manner.

Unlike registers, which are private to each thread, all threads in the same block have access to shared memory.

Amount of shared memory per block:

- 16,384 bytes on the GeForce 9400M,
- 49,152 bytes on the Tesla C2050/C2070,
- 49,152 bytes on the Tesla K20c and the P100.

## constant, global, and cache memory

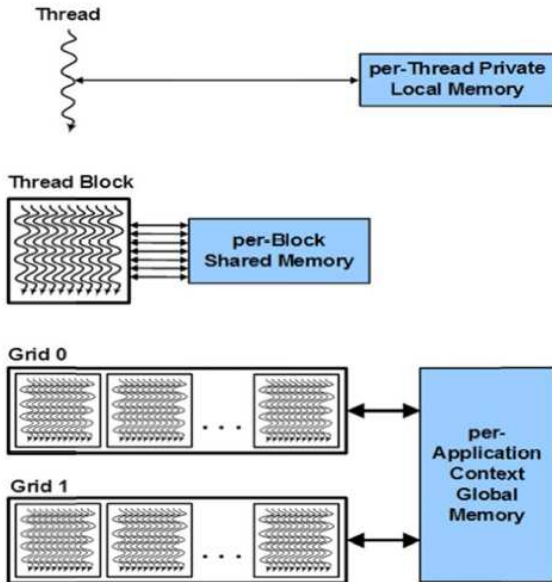
The constant memory supports short-latency, high-bandwidth, read-only access by the device when all threads simultaneously access the same location.

Global memory is similar to RAM on the CPU.

GPU	constant	global	L2 cache
GeForce 9400M	65,536 b	254 Mb	
Tesla C2050	65,536 b	2,687 Mb	786,432 b
Tesla K20C	65,536 b	4,800 Mb	1,310,720 b
Tesla P100	65,536 b	16,276 Mb	4,194,304 b



# a quick refresher



copied from the NVIDIA Whitepaper on Kepler GK110

# Device Memories and Matrix Multiplication

## 1 Device Memories

- global, constant, and shared memories
- **CUDA variable type qualifiers**

## 2 Matrix Multiplication

- an application of tiling
- running `matrixMul` in the GPU Computing SDK
- the kernel of `matrixMul`

# variables in memory, scope, and lifetime

Each variable is stored in a particular type of memory, has a scope and a lifetime.

Scope is the range of threads that can access the variable.

- If the scope of a variable is a single thread, then a private version of that variable exists for every single thread.
- Each thread can access only its private version of the variable.

Lifetime specifies the portion of the duration of the program execution when the variable is available for use.

- If a variable is declared in the kernel function body, then that variable is available for use only by the code of the kernel.
- If the kernel is invoked several times, then the contents of that variable will not be maintained across these invocations.

# CUDA variable type qualifiers

We distinguish between five different variable declarations, based on their memory location, scope, and lifetime.

variable declaration	memory	scope	lifetime
atomic variables $\neq$ arrays	register	thread	kernel
array variables	local	thread	kernel
<code>__device__ __shared__ int v</code>	shared	block	kernel
<code>__device__ int v</code>	global	grid	program
<code>__device__ __constant__ int v</code>	constant	grid	program

The `__device__` in front of `__shared__` is optional.

# Device Memories and Matrix Multiplication

## 1 Device Memories

- global, constant, and shared memories
- CUDA variable type qualifiers

## 2 Matrix Multiplication

- an application of tiling
- running `matrixMul` in the GPU Computing SDK
- the kernel of `matrixMul`

# the CGMA ratio

In our simple matrix-matrix multiplication  $C = A \cdot B$ , we have the statement

```
C[i] += (* (pA++)) * (*pB);
```

where

- $C$  is a float array; and
- $pA$  and  $pB$  are pointers to elements in a float array.

For the statement above, the CGMA ratio is 2/3:

- for one addition and one multiplication,
- we have three memory accesses.

# an application of tiling

For  $A \in \mathbb{R}^{n \times m}$  and  $B \in \mathbb{R}^{m \times p}$ , the product  $C = A \cdot B \in \mathbb{R}^{n \times p}$ .

Assume that  $n$ ,  $m$ , and  $p$  are multiples of some  $w$ , e.g.:  $w = 8$ .

We compute  $C$  in tiles of size  $w \times w$ :

- Every block computes one tile of  $C$ .
- All threads in one block operate on submatrices:


$$C_{i,j} = \sum_{k=1}^{m/w} A_{i,k} \cdot B_{k,j}.$$

- The submatrices  $A_{i,k}$  and  $B_{k,j}$  are loaded from global memory into shared memory of the block.

# matrix multiplication with shared memory

$$C_{i,j} = \sum_{k=1}^{m/w} A_{i,k} \cdot B_{k,j}$$


Diagram illustrating matrix  $B$  with dimensions  $m$  (rows) and  $j$  (columns). A vertical strip of width  $w$  is highlighted, representing the columns  $k$  used in the summation.





# Device Memories and Matrix Multiplication

## 1 Device Memories

- global, constant, and shared memories
- CUDA variable type qualifiers

## 2 Matrix Multiplication

- an application of tiling
- **running `matrixMul` in the GPU Computing SDK**
- the kernel of `matrixMul`

# the `matrixMul` in the GPU Computing SDK

The matrix-matrix multiplication is explained in great detail in the CUDA programming guide.

One of the examples in the GPU Computing SDK is `matrixMul`.

We run it on the GeForce 9400M, the Tesla C2050/C2070, the Tesla K20c (on `kepler`), and P100 (on `pascal`).

# on the GeForce 9400M

```
/Developer/GPU Computing/C/bin/darwin/release $ ./matrixMul  
[matrixMul] starting...
```

```
[ matrixMul ]  
./matrixMul  
Starting (CUDA and CUBLAS tests)...
```

```
Device 0: "GeForce 9400M" with Compute 1.1 capability
```

```
Using Matrix Sizes: A(160 x 320), B(160 x 320), C(160 x 320)
```

```
Runing Kernels...
```

```
> CUBLAS          7.2791 GFlop/s, Time = 0.00225 s, Size = 16384000 Ops
```

```
> CUDA matrixMul 5.4918 GFlop/s, Time = 0.00298 s, Size = 16384000 Ops,  
NumDevsUsed = 1, Workgroup = 256
```

```
Comparing GPU results with Host computation...
```

```
Comparing CUBLAS & Host results  
CUBLAS compares OK
```

```
Comparing CUDA matrixMul & Host results  
CUDA matrixMul compares OK
```

```
[matrixMul] test results...  
PASSED
```

# on the Tesla C2050/C2070

```
/usr/local/cuda/sdk/C/bin/linux/release jan$ ./matrixMul
```

```
[matrixMul] starting...
```

```
[ matrixMul ]
```

```
./matrixMul Starting (CUDA and CUBLAS tests)...
```

```
Device 0: "Tesla C2050 / C2070" with Compute 2.0 capability
```

```
Using Matrix Sizes: A(640 x 960), B(640 x 640), C(640 x 960)
```

```
Runing Kernels...
```

```
> CUBLAS          Throughput = 424.8840 GFlop/s, Time = 0.00185 s, \  
Size = 786432000 Ops
```

```
> CUDA matrixMul Throughput = 186.7684 GFlop/s, Time = 0.00421 s, \  
Size = 786432000 Ops, NumDevsUsed = 1, Workgroup = 1024
```

```
Comparing GPU results with Host computation...
```

```
Comparing CUBLAS & Host results
```

```
CUBLAS compares OK
```

```
Comparing CUDA matrixMul & Host results
```

```
CUDA matrixMul compares OK
```

```
[matrixMul] test results...
```

```
PASSED
```

# on the K20c

```
$ /usr/local/cuda/samples/0_Simple/matrixMul/matrixMul  
[Matrix Multiply Using CUDA] - Starting...
```

```
GPU Device 0: "Tesla K20c" with compute capability 3.5
```

```
MatrixA(320,320), MatrixB(640,320)  
Computing result using CUDA Kernel...  
done
```

```
Performance= 246.13 GFlop/s, Time= 0.533 msec, Size= 131072000 Ops,  
WorkgroupSize= 1024 threads/block  
Checking computed result for correctness: Result = PASS
```

```
Note: For peak performance, please refer to the matrixMulCUBLAS \  
example.  
$
```

The theoretical peak performance of the K20c is 1.17 TFlops double precision, and 3.52 TFlops single precision.

The matrices that are multiplied have single float as type.

# going for peak performance with CUBLAS

```
$ /usr/local/cuda/samples/0_Simple/matrixMulCUBLAS/matrixMulCUBLAS
[Matrix Multiply CUBLAS] - Starting...
/usr/bin/nvidia-modprobe: unrecognized option: "-u"

GPU Device 0: "Tesla K20c" with compute capability 3.5

MatrixA(320,640), MatrixB(320,640), MatrixC(320,640)
Computing result using CUBLAS...done.
Performance= 1171.83 GFlop/s, Time= 0.112 msec, Size= 131072000 Ops
Computing result using host CPU...done.
Comparing CUBLAS Matrix Multiply with CPU results: PASS
$
```

The theoretical peak performance of the K20c is 1.17 TFlops double precision, and 3.52 TFlops single precision.

The matrices that are multiplied have single float as type.

# on the P100

```
$ /usr/local/cuda/samples/0_Simple/matrixMul/matrixMul
[Matrix Multiply Using CUDA] - Starting...
GPU Device 0: "Tesla P100-PCIE-16GB" with compute capability 6.0

MatrixA(320,320), MatrixB(640,320)
Computing result using CUDA Kernel...
done
Performance= 1909.26 GFlop/s, Time= 0.069 msec, Size= 131072000 Ops
WorkgroupSize= 1024 threads/block
Checking computed result for correctness: Result = PASS

NOTE: The CUDA Samples are not meant for performance measurements.
Results may vary when GPU Boost is enabled.

$
```

The theoretical peak performance (with GPU Boost):  
18.7 TFlops (half), 9.3 TFlops (single), 4.7 TFlops (double).

# running CUBLAS on P100

```
$ /usr/local/cuda/samples/0_Simple/matrixMulCUBLAS/matrixMulCUBLAS
[Matrix Multiply CUBLAS] - Starting...
GPU Device 0: "Tesla P100-PCI-E-16GB" with compute capability 6.0

MatrixA(640,480), MatrixB(480,320), MatrixC(640,320)
Computing result using CUBLAS...done.
Performance= 3089.82 GFlop/s, Time= 0.064 msec, Size= 196608000 Ops
Computing result using host CPU...done.
Comparing CUBLAS Matrix Multiply with CPU results: PASS

NOTE: The CUDA Samples are not meant for performance measurements.
Results may vary when GPU Boost is enabled.
$
```

A second run gave the following:

```
Performance= 3106.43 GFlop/s, Time= 0.063 msec, Size= 196608000 Ops
```

For single floats, the theoretical peak performance is 9.3 TFlops.



# Device Memories and Matrix Multiplication

## 1 Device Memories

- global, constant, and shared memories
- CUDA variable type qualifiers

## 2 Matrix Multiplication

- an application of tiling
- running `matrixMul` in the GPU Computing SDK
- the kernel of `matrixMul`

## the kernel of `matrixMul`

```
template <int BLOCK_SIZE> __global__ void
matrixMul( float* C, float* A, float* B, int wA, int wB)
{
    int bx = blockIdx.x;    // Block index
    int by = blockIdx.y;
    int tx = threadIdx.x;   // Thread index
    int ty = threadIdx.y;
    // Index of the first sub-matrix of A processed by the block
    int aBegin = wA * BLOCK_SIZE * by;
    // Index of the last sub-matrix of A processed by the block
    int aEnd   = aBegin + wA - 1;
    // Step size used to iterate through the sub-matrices of A
    int aStep  = BLOCK_SIZE;
    // Index of the first sub-matrix of B processed by the block
    int bBegin = BLOCK_SIZE * bx;
    // Step size used to iterate through the sub-matrices of B
    int bStep  = BLOCK_SIZE * wB;
```

## the submatrices

```
// Csub is used to store the element of the block sub-matrix
// that is computed by the thread
float Csub = 0;

// Loop over all the sub-matrices of A and B
// required to compute the block sub-matrix
for (int a = aBegin, b = bBegin;
     a <= aEnd;
     a += aStep, b += bStep) {

    // Declaration of the shared memory array As used to
    // store the sub-matrix of A
    __shared__ float As[BLOCK_SIZE][BLOCK_SIZE];

    // Declaration of the shared memory array Bs used to
    // store the sub-matrix of B
    __shared__ float Bs[BLOCK_SIZE][BLOCK_SIZE];
```

# loading and multiplying

```
// Load the matrices from device memory
// to shared memory; each thread loads
// one element of each matrix
AS(ty, tx) = A[a + wA * ty + tx];
BS(ty, tx) = B[b + wB * ty + tx];

// Synchronize to make sure the matrices are loaded
__syncthreads();

// Multiply the two matrices together;
// each thread computes one element
// of the block sub-matrix
#pragma unroll
for (int k = 0; k < BLOCK_SIZE; ++k)
    Csub += AS(ty, k) * BS(k, tx);

// Synchronize to make sure that the preceding
// computation is done before loading two new
// sub-matrices of A and B in the next iteration
__syncthreads();
}
```

# the end of the kernel

```
// Write the block sub-matrix to device memory;  
// each thread writes one element  
int c = wB * BLOCK_SIZE * by + BLOCK_SIZE * bx;  
C[c + wB * ty + tx] = Csub;  
}
```

The emphasis in this lecture is on

- 1 the use of device memories; and
- 2 data organization (tiling) and transfer.

In the next lecture we will come back to this code,  
and cover thread scheduling

- 1 the use of `blockIdx`; and
- 2 thread synchronization.

## summary and exercises

Vasily Volkov and James W. Demmel: **Benchmarking GPUs to tune dense linear algebra**. In *Proceedings of the 2008 ACM/IEEE conference on Supercomputing*. IEEE Press, 2008. Article No. 31.

We covered more of chapter 3 in the book of Kirk & Hwu, and also several concepts explained in chapter 5.

- 1 Compile the `matrixMul` of the GPU Computing SDK on your laptop and desktop and run the program.
- 2 Consider the matrix multiplication code of last lecture and compute the CGMA ratio.
- 3 Adjust the code for matrix multiplication we discussed last time to use shared memory.