

Thread Organization and Matrix Multiplication

1 Thread Organization

- grids, blocks, and threads
- using `threadIdx` and `blockIdx`
- setting the execution configuration parameters

2 Matrix Matrix Multiplication

- accessing submatrices with thread identifiers
- CUDA code for thread organization
- thread synchronization
- revisiting the kernel of `matrixMul`

MCS 572 Lecture 33
Introduction to Supercomputing
Jan Verschelde, 7 November 2016

Thread Organization and Matrix Multiplication

1 Thread Organization

- grids, blocks, and threads
- using `threadIdx` and `blockIdx`
- setting the execution configuration parameters

2 Matrix Matrix Multiplication

- accessing submatrices with thread identifiers
- CUDA code for thread organization
- thread synchronization
- revisiting the kernel of `matrixMul`

grids, blocks, and threads

The code that runs on the GPU is defined in a function, the kernel.

A kernel launch

- creates a grid of blocks, and
- each block has one or more threads.

The organization of the grids and blocks can be 1D, 2D, or 3D.

During the running of the kernel:


- Threads in the same block are executed simultaneously.
- Blocks are scheduled by the streaming multiprocessors.

The NVIDIA Tesla C2050 has 14 streaming multiprocessors and threads are executed in groups of 32 (the warp size).

This implies: $14 \times 32 = 448$ threads can run simultaneously.

For the K20c the numbers are respectively 13, 192, and 2496; and for the P100, we have 56, 64, and 3584.

a scalable programming model



Thread Organization and Matrix Multiplication

1 Thread Organization

- grids, blocks, and threads
- **using** `threadIdx` **and** `blockIdx`
- setting the execution configuration parameters

2 Matrix Matrix Multiplication

- accessing submatrices with thread identifiers
- CUDA code for thread organization
- thread synchronization
- revisiting the kernel of `matrixMul`

identifying threads

All threads execute the same code, defined by the kernel.

The builtin variable `threadIdx`

- identifies every thread in a block uniquely; and
- defines the data processed by the thread.

The builtin variable `blockDim` holds the number of threads in a block.


In a one dimensional organization, we use only `threadIdx.x` and `blockDim.x`. For 2D and 3D, the other components

- `threadIdx.y` belongs to the range `0..blockDim.y`;
- `threadIdx.z` belongs to the range `0..blockDim.z`.

data for each thread

The grid consists of N blocks, with $\text{blockIdx.x} \in \{0, N - 1\}$.

Within each block, $\text{threadIdx.x} \in \{0, \text{blockDim.x} - 1\}$.



```
int threadId = blockIdx.x *  
    blockDim.x + threadIdx.x  
...  
float x = input[threadID]  
float y = f(x)  
output[threadID] = y  
...
```

Thread Organization and Matrix Multiplication

1 Thread Organization

- grids, blocks, and threads
- using `threadIdx` and `blockIdx`
- setting the execution configuration parameters

2 Matrix Matrix Multiplication

- accessing submatrices with thread identifiers
- CUDA code for thread organization
- thread synchronization
- revisiting the kernel of `matrixMul`

setting the execution configuration parameters

Suppose the kernel is defined by the function F with input arguments x and output arguments y , then

```
dim3 dimGrid(128, 1, 1);  
dim3 dimBlock(32, 1, 1);  
F<<<dimGrid, dimBlock>>>(x, y);
```

launches a grid of 128 blocks. The grid is a one dimensional array. Each block in the grid is also one dimensional and has 32 threads.

multidimensional thread organization

Limitations of the Tesla C2050/C2070:

- Maximum number of threads per block: 1,024.
- Maximum sizes of each dimension of a block: $1,024 \times 1,024 \times 64$.
Because 1,024 is the upper limit for the number of threads in a block, the largest square 2D block is 32×32 , as $32^2 = 1,024$.
- Maximum sizes of each dimension of a grid:
 $65,535 \times 65,535 \times 65,535$.
65,535 is the upper limit for the builtin variables `gridDim.x`, `gridDim.y`, and `gridDim.z`.

Limitations of the K20c and the P100:

- Maximum number of threads per block: 1,024.
- Maximum dimension size of a thread block: $1,024 \times 1,024 \times 64$.
- Maximum dimension size of a grid size:
 $2,147,483,647 \times 65,535 \times 65,535$

a 3D example

Suppose the function F defines the kernel,
with argument x , then

```
dim3 dimGrid(3, 2, 4);  
dim3 dimBlock(5, 6, 2);  
F<<<dimGrid, dimBlock>>>(x);
```

launches a grid with

- $3 \times 2 \times 4$ blocks; and
- each block has $5 \times 6 \times 2$ threads.

Thread Organization and Matrix Multiplication

1 Thread Organization


- grids, blocks, and threads
- using `threadIdx` and `blockIdx`
- setting the execution configuration parameters

2 Matrix Matrix Multiplication

- accessing submatrices with thread identifiers
- CUDA code for thread organization
- thread synchronization
- revisiting the kernel of `matrixMul`



submatrices

Consider a grid of dimension $2 \times 2 \times 1$
to store a 4-by-4 matrix in tiles of dimensions $2 \times 2 \times 1$:




mapping threads to entries in the matrix

A kernel launch with a grid of dimensions $2 \times 2 \times 1$ where each block has dimensions $2 \times 2 \times 1$ creates 16 threads.



linear address calculation

A kernel launch with a grid of dimensions $2 \times 2 \times 1$
where each block has dimensions $2 \times 2 \times 1$ creates 16 threads.



```
x[0][0][0][0][0][0] = 0
x[0][0][0][0][1][0] = 1
x[0][0][0][1][0][0] = 2
x[0][0][0][1][1][0] = 3
x[0][1][0][0][0][0] = 4
x[0][1][0][0][1][0] = 5
x[0][1][0][1][0][0] = 6
x[0][1][0][1][1][0] = 7
x[1][0][0][0][0][0] = 8
x[1][0][0][0][1][0] = 9
x[1][0][0][1][0][0] = 10
x[1][0][0][1][1][0] = 11
x[1][1][0][0][0][0] = 12
x[1][1][0][0][1][0] = 13
x[1][1][0][1][0][0] = 14
x[1][1][0][1][1][0] = 15
```

Thread Organization and Matrix Multiplication

1 Thread Organization

- grids, blocks, and threads
- using `threadIdx` and `blockIdx`
- setting the execution configuration parameters

2 Matrix Matrix Multiplication

- accessing submatrices with thread identifiers
- **CUDA code for thread organization**
- thread synchronization
- revisiting the kernel of `matrixMul`

the main program

```
int main ( int argc, char* argv[] )
{
    const int xb = 2; /* gridDim.x */
    const int yb = 2; /* gridDim.y */
    const int zb = 1; /* gridDim.z */
    const int xt = 2; /* blockDim.x */
    const int yt = 2; /* blockDim.y */
    const int zt = 1; /* blockDim.z */
    const int n = xb*yb*zb*xt*yt*zt;

    printf("allocating array of length %d...\n",n);

    /* allocating and initializing on the host */

    int *xhost = (int*)calloc(n,sizeof(int));
    for(int i=0; i<n; i++) xhost[i] = -1.0;
```

copy to device and kernel launch

```
int *xdevice;
size_t sx = n*sizeof(int);
cudaMalloc((void**)&xdevice, sx);
cudaMemcpy(xdevice, xhost, sx, cudaMemcpyHostToDevice);

/* set the execution configuration for the kernel */

dim3 dimGrid(xb, yb, zb);
dim3 dimBlock(xt, yt, zt);
matrixFill<<<dimGrid, dimBlock>>>(xdevice);
```

the kernel definition

```
__global__ void matrixFill ( int *x )  
/*  
 * Fills the matrix using blockIdx and threadIdx. */  
{  
    int bx = blockIdx.x;  
    int by = blockIdx.y;  
    int tx = threadIdx.x;  
    int ty = threadIdx.y;  
    int row = by*blockDim.y + ty;  
    int col = bx*blockDim.x + tx;  
    int dim = gridDim.x*blockDim.x;  
    int i = row*dim + col;  
    x[i] = i;  
}
```

copying to host and writing the result

```
/* copy data from device to host */
cudaMemcpy(xhost,xdevice,sx,cudaMemcpyDeviceToHost);
cudaFree(xdevice);

int *p = xhost;
for(int i1=0; i1 < xb; i1++)
    for(int i2=0; i2 < yb; i2++)
        for(int i3=0; i3 < zb; i3++)
            for(int i4=0; i4 < xt; i4++)
                for(int i5=0; i5 < yt; i5++)
                    for(int i6=0; i6 < zt; i6++)
                        printf("x[%d][%d][%d][%d][%d][%d] = %d\n",
                               i1,i2,i3,i4,i5,i6,*p++));

return 0;
}
```

Thread Organization and Matrix Multiplication

1 Thread Organization

- grids, blocks, and threads
- using `threadIdx` and `blockIdx`
- setting the execution configuration parameters

2 Matrix Matrix Multiplication

- accessing submatrices with thread identifiers
- CUDA code for thread organization
- **thread synchronization**
- revisiting the kernel of `matrixMul`

thread synchronization

In a block all threads run independently.

CUDA allows threads in the same block to coordinate their activities using a barrier synchronization function:


```
__syncthreads();
```

The thread executing `__syncthreads()` will be held at the calling location in the code until every thread in the block reaches the location.

Placing a `__syncthreads()` ensures that all threads in a block have completed a task before moving on.

applied to matrix multiplication with shared memory

$$C_{i,j} = \sum_{k=1}^{m/w} A_{i,k} \cdot B_{k,j}$$

application of `__syncthreads()`

With tiled matrix matrix multiplication using shared memory, all threads in the block collaborate to copy the tiles $A_{i,k}$ and $B_{k,j}$ from global memory to shared memory.

→ Before the calculation of the inner products, all threads must finish their copy statement: they all execute the `__syncthreads()`.

Every thread computes one inner product.

→ Before moving on to the next tile, all threads must finish, therefore, they all execute the `__syncthreads()` after computing their inner product and moving on to the next phase.

Thread Organization and Matrix Multiplication

1 Thread Organization

- grids, blocks, and threads
- using `threadIdx` and `blockIdx`
- setting the execution configuration parameters

2 Matrix Matrix Multiplication

- accessing submatrices with thread identifiers
- CUDA code for thread organization
- thread synchronization
- revisiting the kernel of `matrixMul`

the kernel of matrixMul

```
template <int BLOCK_SIZE> __global__ void
matrixMul( float* C, float* A, float* B, int wA, int wB)
{
    int bx = blockIdx.x;    // Block index
    int by = blockIdx.y;
    int tx = threadIdx.x;   // Thread index
    int ty = threadIdx.y;
    // Index of the first sub-matrix of A processed by the block
    int aBegin = wA * BLOCK_SIZE * by;
    // Index of the last sub-matrix of A processed by the block
    int aEnd   = aBegin + wA - 1;
    // Step size used to iterate through the sub-matrices of A
    int aStep  = BLOCK_SIZE;
    // Index of the first sub-matrix of B processed by the block
    int bBegin = BLOCK_SIZE * bx;
    // Step size used to iterate through the sub-matrices of B
    int bStep  = BLOCK_SIZE * wB;
```

the submatrices

```
// Csub is used to store the element of the block sub-matrix  
// that is computed by the thread  
float Csub = 0;  
  
// Loop over all the sub-matrices of A and B  
// required to compute the block sub-matrix  
for (int a = aBegin, b = bBegin;  
     a <= aEnd;  
     a += aStep, b += bStep) {  
  
    // Declaration of the shared memory array As used to  
    // store the sub-matrix of A  
    __shared__ float As[BLOCK_SIZE][BLOCK_SIZE];  
  
    // Declaration of the shared memory array Bs used to  
    // store the sub-matrix of B  
    __shared__ float Bs[BLOCK_SIZE][BLOCK_SIZE];
```

loading and multiplying

```
// Load the matrices from device memory
// to shared memory; each thread loads
// one element of each matrix
AS(ty, tx) = A[a + wA * ty + tx];
BS(ty, tx) = B[b + wB * ty + tx];

// Synchronize to make sure the matrices are loaded
__syncthreads();

// Multiply the two matrices together;
// each thread computes one element
// of the block sub-matrix
#pragma unroll
for (int k = 0; k < BLOCK_SIZE; ++k)
    Csub += AS(ty, k) * BS(k, tx);

// Synchronize to make sure that the preceding
// computation is done before loading two new
// sub-matrices of A and B in the next iteration
__syncthreads();
}
```

the end of the kernel

```
// Write the block sub-matrix to device memory;  
// each thread writes one element  
int c = wB * BLOCK_SIZE * by + BLOCK_SIZE * bx;  
C[c + wB * ty + tx] = Csub;  
}
```

Recommended reading:

- NVIDIA CUDA Programming Guide.
Available at <http://developer.nvidia.com>.
- Vasily Volkov and James W. Demmel: **Benchmarking GPUs to tune dense linear algebra**. In *Proceedings of the 2008 ACM/IEEE conference on Supercomputing*. IEEE Press, 2008. Article No. 31.

summary and exercises

We covered more of chapter 4 in the book of Kirk & Hwu.

- 1 Find the limitations of the grid and block sizes for the graphics card on your laptop or desktop.
- 2 Extend the simple code with the three dimensional thread organization to a tiled matrix-vector multiplication for numbers generated at random as 0 or 1.